

Introduction to probability theory

Computational modeling of language
Spring 2026

Bureaucracy

- **Discussion board:** Now in place on bCourses.

Bureaucracy

- **Discussion board:** Now in place on bCourses.
- **In case of illness:** Please stay home and get better - we will accommodate!

Last week

- The pumping lemma
- PDAs and context-free languages
- Chomsky vs. Shannon on human language
- On toward numbers
- Pandas, Karen, Terry

Last week

- The pumping lemma
- PDAs and context-free languages
- Chomsky vs. Shannon on human language
- On toward numbers
- Pandas, Karen, Terry

This week

- Introduction to probability theory
- Bayes' rule
- Marr's levels of analysis
- Markov chains

Small groups



Trade notes on **last week's material**,
and **this week's readings**, to generate
comments and/or questions.



Last week



- The pumping lemma
- PDAs and context-free languages
- Chomsky vs. Shannon on human language
- On toward numbers
- Pandas, Karen, Terry

This week

- Introduction to probability theory
- Bayes' rule
- Marr's levels of analysis
- Markov chains










Probability theory


a consistent framework for the quantification
and manipulation of **uncertainty**

Bishop 2006






- Random variable: a variable whose possible values reflect outcomes of a random phenomenon.




- Random variable: a variable whose possible values reflect outcomes of a random phenomenon.
- E.g. B : the color of a selected box (r or b)




$$p(B=r) = .4$$
$$p(B=b) = .6$$

- Random variable: a variable whose possible values reflect outcomes of a random phenomenon.
- E.g. B : the color of a selected box (r or b)



$$p(B=r) = .4$$
$$p(B=b) = .6$$

- Random variable: a variable whose possible values reflect outcomes of a random phenomenon.
- E.g. B : the color of a selected box (r or b)
- E.g. F : selected fruit (a or o)




$$p(B=r) = .4$$

$$p(B=b) = .6$$

$$p(F=a|B=r) =$$

$$p(F=o|B=r) =$$

- Random variable: a variable whose possible values reflect outcomes of a random phenomenon.
- E.g. B : the color of a selected box (r or b)
- E.g. F : selected fruit (a or o)



$$p(B=r) = .4$$

$$p(B=b) = .6$$

$$p(F=a|B=r) = 2/8$$

$$p(F=o|B=r) = 6/8$$


- Random variable: a variable whose possible values reflect outcomes of a random phenomenon.
- E.g. B : the color of a selected box (r or b)
- E.g. F : selected fruit (a or o)

Interpretation of probability

Frequentist vs. Bayesian

Interpretation of probability

Frequentist: Probability measures **proportion of outcomes in repeated trials**.



$$p(F=a|B=r) = 2/8$$

$$p(F=o|B=r) = 6/8$$

Interpretation of probability

Frequentist: Probability measures **proportion of outcomes in repeated trials**.

Interpretation of probability

Frequentist: Probability measures **proportion of outcomes in repeated trials**.

- $p(\textit{Weather}=\textit{rain}|\textit{Date}=\textit{Feb 11})$ is (the limit of) the proportion of February 11s in the historical record during which it rained.

Interpretation of probability

Frequentist: Probability measures **proportion of outcomes in repeated trials**.

- $p(\textit{Weather}=\textit{rain}|\textit{Date}=\textit{Feb 11})$ is (the limit of) the proportion of February 11s in the historical record during which it rained.
- $p(\textit{Weather}=\textit{rain}|\textit{Date}=\textit{Feb 11 2026})$ is strictly meaningless – or trivially either 0 or 1, depending on whether it rains tomorrow or not.

Interpretation of probability

Bayesian: Probability measures **degree of belief**.

Interpretation of probability

Bayesian: Probability measures **degree of belief**.

- Here, $p(\textit{Weather}=\textit{rain}|\textit{Date}=\textit{Feb 11 2026})$ is meaningful.

Interpretation of probability

Bayesian: Probability measures **degree of belief**.

- Here, $p(\textit{Weather}=\textit{rain}|\textit{Date}=\textit{Feb 11 2026})$ is meaningful.
- It measures the **degree of belief** that it will rain, given that the date is Feb. 11 2026.

2 random variables: X , Y

- X can take on values $x_1, x_2, x_3, \dots$
- Y can take on values $y_1, y_2, y_3, \dots$

2 random variables: X , Y

- X can take on values $x_1, x_2, x_3, \dots$
- Y can take on values $y_1, y_2, y_3, \dots$
- And we can **count how often** particular **pairs of values** x_i, y_j occur. That count is n_{ij} .

2 random variables: X , Y

y_j			n_{ij}	

x_i

- X can take on values $x_1, x_2, x_3, \dots$
- Y can take on values $y_1, y_2, y_3, \dots$
- And we can **count how often** particular **pairs of values** x_i, y_j occur. That count is n_{ij} .

n_{ij} = count for x_i, y_j

y_j			n_{ij}	

x_i

n_{ij} = count for x_i, y_j


c_i = total count in column i

y_j			n_{ij}	
			x_i	

n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j




n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells




n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells

$p(X=x_i, Y=y_j) = ?$




n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells

$$p(X=x_i, Y=y_j) = n_{ij} / N$$




n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells

$p(X=x_i) = ?$




n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells

$$p(X=x_i) = c_i/N$$




n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells

$$p(X=x_i) = c_i/N = (n_{i1} + n_{i2} + n_{i3}) / N$$




n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells



$$p(X=x_i) = c_i/N = (n_{i1} + n_{i2} + n_{i3}) / N$$


$$= (n_{i1}/N) + (n_{i2}/N) + (n_{i3}/N)$$

n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells



$$p(X=x_i) = c_i/N = (n_{i1} + n_{i2} + n_{i3}) / N$$

$$= (n_{i1}/N) + (n_{i2}/N) + (n_{i3}/N)$$


$$= p(X=x_i, Y=y_1) + p(X=x_i, Y=y_2) + p(X=x_i, Y=y_3)$$

n_{ij} = count for x_i, y_j

c_i = total count in column i

r_j = total count in row j

N = total count in all cells



Sum rule:
$$p(X = x_i) = \sum_{j=1}^L p(X = x_i, Y = y_j)$$

The Rules of Probability

sum rule

$$p(X) = \sum_Y p(X, Y)$$

product rule

$$p(X, Y) = p(Y|X)p(X).$$

Bayes' rule

$$p(A|B) = \frac{p(B|A) p(A)}{p(B)}$$

Bayes' rule

$$\text{posterior}$$
$$p(A|B) = \frac{p(B|A) p(A)}{p(B)}$$

Bayes' rule

$$\underbrace{p(A|B)}_{\text{posterior}} = \frac{\underbrace{p(B|A)}_{\text{likelihood}} \underbrace{p(A)}_{\text{prior}}}{p(B)}$$

Bayes' rule

$$\underbrace{p(A|B)}_{\text{posterior}} = \frac{\underbrace{p(B|A)}_{\text{likelihood}} \underbrace{p(A)}_{\text{prior}}}{\underbrace{p(B)}_{\text{normalizing constant}}}$$

Bayes' rule

$$\begin{array}{c} \text{posterior} \\ \hline p(A|B) \end{array} = \frac{\begin{array}{c} \text{likelihood} \\ \hline p(B|A) \end{array} \begin{array}{c} \text{prior} \\ \hline p(A) \end{array}}{\begin{array}{c} \hline \\ \text{normalizing} \\ \text{constant} \end{array} p(B)}$$

$$\begin{array}{c} \text{posterior} \\ \hline p(A|B) \end{array} \propto \begin{array}{c} \text{likelihood} \\ \hline p(B|A) \end{array} \begin{array}{c} \text{prior} \\ \hline p(A) \end{array}$$

Bayes' rule

$$\frac{\text{posterior}}{p(H|d)} \propto \frac{\text{likelihood}}{p(d|H)} \frac{\text{prior}}{p(H)}$$

- H : hypothesis
- d : data

Bayes' Rule

A Tutorial Introduction to Bayesian Analysis

James V Stone

$$\frac{p(x|\theta)p(\theta)}{p(x)}$$

$$p(\theta|x)$$

$$\begin{aligned} p(\theta|x_1) &= \frac{p(x_1|\theta)p(\theta)}{p(x_1)} \\ &= \theta^1(1-\theta)^{1-1} \\ &= \theta \end{aligned}$$

X/θ	θ_1	θ_2	θ_3	θ_4	θ_5	θ_6	θ_7	θ_8	θ_9	θ_{10}
x_4	0	0	1	0	3	5	10	7	7	4
x_3	0	1	1	10	16	11	12	7	8	5
x_2	3	5	8	9	14	19	3	3	0	0
x_1	8	9	9	5	4	4	1	0	0	0
Sum	11	19	24	37	27	26	14	7	15	7

I am
My Lord
Your Lordship's
most obedient
humble servant
J. Bayes.

It is remarkable that a science which began with the consideration of games of chance should have become the most important object of human knowledge.

Pierre Simon Laplace, 1812.

“... we balance probabilities and choose the most likely. It is the scientific use of the imagination ... ”

Sherlock Holmes, The Hound of the Baskervilles.

AC Doyle, 1901.

?

“... we balance probabilities and choose the most likely. It is the scientific use of the imagination .. ”

Sherlock Holmes, The Hound of the Baskervilles.

AC Doyle, 1901.



(a) Bayes



(b) Laplace

Figure 1.1.: The fathers of Bayes' rule. a) Thomas Bayes (c. 1701-1761).
b) Pierre-Simon Laplace (1749-1827).

J. V. Stone 2013







Figure 1.5.: Thomas Bayes trying to make sense of a London accent, which removes the *h* sound from the word *handle*, so the phrase *fork handles* is pronounced *fork 'andles*, and therefore sounds like *four candles* (see Fork Handles YouTube clip by The Two Ronnies).



$$\frac{\text{posterior}}{p(H|d)} \propto \frac{\text{likelihood}}{p(d|H)} \frac{\text{prior}}{p(H)}$$




$$\frac{\text{posterior}}{p(H|d)} \propto \frac{\text{likelihood}}{p(d|H)} \frac{\text{prior}}{p(H)}$$



Assume:

$$p(d|H1) = p(d|H2)$$

$$\underbrace{p(H|d)}_{\text{posterior}} \propto \underbrace{p(d|H)}_{\text{likelihood}} \underbrace{p(H)}_{\text{prior}}$$




Assume:

$$p(d|H1) = p(d|H2)$$

$$\underbrace{p(H|d)}_{\text{posterior}} \propto \underbrace{p(d|H)}_{\text{likelihood}} \underbrace{p(H)}_{\text{prior}}$$

Uttered in a candle store:
 $p(H1) = .9$; $p(H2) = .1$



Seeing is not a direct apprehension of reality, as we often like to pretend. Quite the contrary: **seeing is inference from incomplete information...**

E.T. Jaynes 2003

Seeing is not a direct apprehension of reality, as we often like to pretend. Quite the contrary: **seeing is inference from incomplete information...**

E.T. Jaynes 2003

See also Helmholtz on
[perception as unconscious inference](#).

DID THE SUN JUST EXPLODE?

(IT'S NIGHT, SO WE'RE NOT SURE.)

THIS NEUTRINO DETECTOR MEASURES
WHETHER THE SUN HAS GONE NOVA.

THEN, IT ROLLS TWO DICE. IF THEY
BOTH COME UP SIX, IT LIES TO US.
OTHERWISE, IT TELLS THE TRUTH.

LET'S TRY.

DETECTOR! HAS THE
SUN GONE NOVA?

ROLL
YES.



DID THE SUN JUST EXPLODE? (IT'S NIGHT, SO WE'RE NOT SURE.)

THIS NEUTRINO DETECTOR MEASURES
WHETHER THE SUN HAS GONE NOVA.

THEN, IT ROLLS TWO DICE. IF THEY
BOTH COME UP SIX, IT LIES TO US.
OTHERWISE, IT TELLS THE TRUTH.

LET'S TRY.

DETECTOR! HAS THE
SUN GONE NOVA?

ROLL

YES.



FREQUENTIST STATISTICIAN:

THE PROBABILITY OF THIS RESULT
HAPPENING BY CHANCE IS $\frac{1}{36} = 0.027$.

SINCE $p < 0.05$, I CONCLUDE
THAT THE SUN HAS EXPLODED.



DID THE SUN JUST EXPLODE? (IT'S NIGHT, SO WE'RE NOT SURE.)

THIS NEUTRINO DETECTOR MEASURES
WHETHER THE SUN HAS GONE NOVA.

THEN, IT ROLLS TWO DICE. IF THEY
BOTH COME UP SIX, IT LIES TO US.
OTHERWISE, IT TELLS THE TRUTH.

LET'S TRY.

DETECTOR! HAS THE
SUN GONE NOVA?

ROLL
YES.



FREQUENTIST STATISTICIAN:

THE PROBABILITY OF THIS RESULT
HAPPENING BY CHANCE IS $\frac{1}{36} = 0.027$.
SINCE $p < 0.05$, I CONCLUDE
THAT THE SUN HAS EXPLODED.



BAYESIAN STATISTICIAN:


BET YOU \$50
IT HASN'T.



Probabilistic inference, using the ideas we have seen here, is often taken as a **normative standard** against which human performance can be compared.

Marr's (1982) levels of analysis

The **Marr Prize**, named in honor of David Marr, is given annually to the **best student paper** at the Cognitive Science Society conference.



COOPERATIVE MINDS
SOCIAL INTERACTION AND GROUP
DYNAMICS

35th Annual Cognitive
Science Conference
Humboldt-Universität
zu Berlin, Germany
July 31–August 3, 2013

KEYNOTE TALKS

Michael E. Bratman (Stanford University)
Cynthia Breazeal (MIT)
Elizabeth Spelke (Harvard University)

INVITED SYMPOSIA

Joint Action
Language and Gesture Evolution
New Frameworks of Rationality

PROGRAM CO-CHAIRS

Markus Knauff, Michael Pauen, Natalie Sebanz,
and Ipke Wachsmuth



Design by Sebastian Lehart, Munich



cognitivesciencesociety.org

Marr's three levels

The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

Marr's three levels

The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

What is the goal
of the
computation,
why is it
appropriate, and
what is the logic
of the strategy
by which it can
be carried out?

Marr's three levels

The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

What is the goal of the computation, why is it appropriate, and what is the logic of the strategy by which it can be carried out?

How can this computational theory be implemented? In particular, what is the representation for the input and output, and what is the algorithm for the transformation?

Marr's three levels

The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

What is the goal of the computation, why is it appropriate, and what is the logic of the strategy by which it can be carried out?

How can this computational theory be implemented? In particular, what is the representation for the input and output, and what is the algorithm for the transformation?

How can the representation and algorithm be realized physically?

Marr's three levels

The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

The functional goal of the visual system is to determine the shape of objects in the world.

Marr's three levels


The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

The functional goal of the visual system is to determine the shape of objects in the world.



Marr's three levels


The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

The functional goal of the visual system is to determine the shape of objects in the world.



Neurons,
somehow.

Marr's three levels

The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

The goal is to
fly.

Curved wings,
aerodynamics

Feathers, balsa
wood, steel,
other...

Marr's three levels

The three levels at which any machine carrying out an information-processing task must be understood

*Computational
theory*

*Representation and
algorithm*

*Hardware
implementation*

The goal is to
fly.

Curved wings,
aerodynamics

Feathers, balsa
wood, steel,
other...

What about language?

The computational level

Specifies “what the system is up to” – what the **functional goal** of its processing is.

The computational level

Specifies “what the system is up to” – what the **functional goal** of its processing is.

Allows us to ask: what is the **optimal** solution for that goal?

The computational level

Specifies “what the system is up to” – what the **functional goal** of its processing is.

Allows us to ask: what is the **optimal** solution for that goal?

And do humans do it that way – regardless of algorithmic and implementational detail?

The computational level

Marr: “trying to understand perception by studying only neurons is like trying to understand bird flight by studying only feathers. It just cannot be done.”

Probabilistic inference, using the ideas we have seen here, is often taken as a **normative standard** against which human performance can be compared.

Probabilistic inference, using the ideas we have seen here, is often taken as a **normative standard** against which human performance can be compared.

This concerns Marr's **computational** level.